Upload Service

Requirements & Acceptance Criteria

RealityCheck visual-analysis test fixture · Version 1.4

Owner	Platform Engineering	Status	Approved for implementation
Target environment	Production	Document type	Requirements / acceptance criteria
Fixture purpose	UI demonstration and cross-source evaluation	Data	Synthetic, non-production

How to use this fixture
Upload this PDF for a single-document analysis. For the richest comparison view, upload it together with **implementation-notes.txt**, **test-observations.txt**, and **production-config.txt** from the companion ZIP. The companion files intentionally contain conflicts and unverified behavior so the UI can show findings, claims, actions, charts, and provenance.

Important interpretation rule

This requirements document establishes expected behavior. It does not prove that implementation satisfies the requirements or that tests have passed. Keep requirements, implementation descriptions, observations, and configuration values separate when interpreting findings.

System context

The Upload Service accepts user-provided evidence files, stores them locally, creates an asynchronous processing job, and exposes job status through an API. The service is expected to be reliable, observable, and safe when handling malformed or unsupported input.

Requirements

R-001 — Supported file types

The service shall accept plain-text (.txt) and PDF (.pdf) evidence files. Other file types shall be rejected with a clear validation error.

R-002 — Maximum upload size

The maximum accepted file size shall be 80 MB per file in the production environment. The limit must be enforced before document extraction begins.

R-003 — Processing timeout

Document extraction shall have a hard processing timeout of 45 seconds per file. A timeout shall result in a terminal failed job state and a user-readable error.

R-004 — Job status

Every accepted upload shall receive a unique job identifier. The API shall expose queued, processing, completed, and failed states.

R-005 — Retention

Uploaded artifacts shall be retained for no longer than 7 days unless an approved retention exception is recorded.

R-006 — Secrets and credentials

The service shall not log API keys, authorization headers, or document contents. Error messages returned to users shall not include stack traces or local filesystem paths.

R-007 — Evidence traceability

Every finding shown to the user shall link to the source document and the relevant extracted evidence. If a conclusion is inferred rather than directly stated, the UI shall label it as inferred or uncertain.

R-008 — Offline mode

When the analysis provider is configured as offline, the UI shall clearly disclose that analysis is simulated and shall not present simulated output as a live model result.

Acceptance tests

ID	Test	Expected result
AT-001	Valid .txt smaller than 80 MB	Expect HTTP 202 and a job identifier.
AT-002	Valid PDF smaller than 80 MB	Expect HTTP 202 and a job identifier.
AT-003	Upload a .csv file	Expect a clear unsupported-type response; no processing job should be created.
AT-004	Upload an 81 MB file	Expect rejection before extraction begins.
AT-005	PDF extraction exceeds 45 seconds	Expect a terminal failed job state and a user-readable timeout error.
AT-006	Inspect logs using a synthetic API key	Confirm the key and document body are not present.
AT-007	Open a finding and follow its source reference	Confirm it reaches the correct evidence excerpt.
AT-008	Set ANALYSIS_PROVIDER=offline	Confirm the UI displays an explicit offline/simulated label.

Known open questions

• Q-001: Does “80 MB” mean decimal megabytes (80,000,000 bytes) or mebibytes (83,886,080 bytes)?

• Q-002: Should a retention exception require an expiry date and an approver?

• Q-003: Should a timed-out extraction be retried automatically, or should the user explicitly retry?

• Q-004: Should the upload limit be configurable per environment, and where should the effective runtime value be displayed?

Fixture limitations: synthetic data only. The companion implementation, configuration, and observation files are intentionally not fully aligned. Findings from the UI should be reviewed as candidates, not treated as verified production defects.